MACHINE LEARNING (FOR PHYSICISTS)

Luc Le Pottier University of Michigan 1/30/2020

(Partially based on materials from Eric Eaton)

What is ML?

"the study of algorithms that improve their performance *P* at task *T* with experience *E*."

- A well-defined machine learning task is one for which P,
 T, and E are defined.
- Examples?
 - best method of navigating undersea

Why?

ML is useful when:

- We have large datasets (particle physics)
- We want to create things that can do things we can't (robots)
- We want to customize models (recommendation systems)

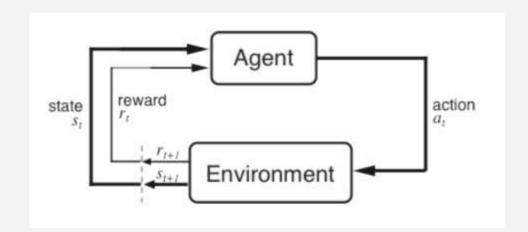
It is *not* always useful (but this is frequently disregarded)

Types of Learning

- Supervised Learning
 - Given data samples X and their labels Y
 - Regression: learn a function f(x) to predict y given x
 - Classification: regression, but y is categorical
- Unsupervised Learning
 - Given data samples X only
 - Want to determine hidden structure in some way

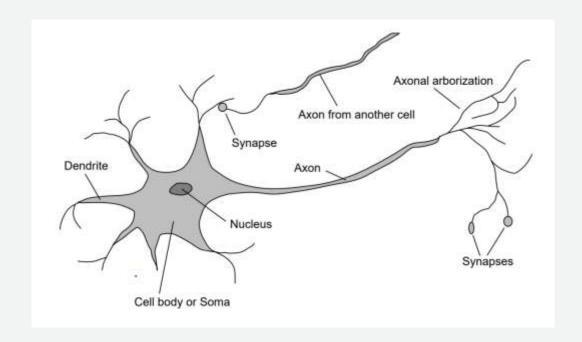
Types of Learning

- Reinforcement Learning
 - Given a sequence of states +
 actions with delayed rewards
 - Determine a mapping from states
 to actions which describes the
 correct action given a state
 - Playing video games, etc.



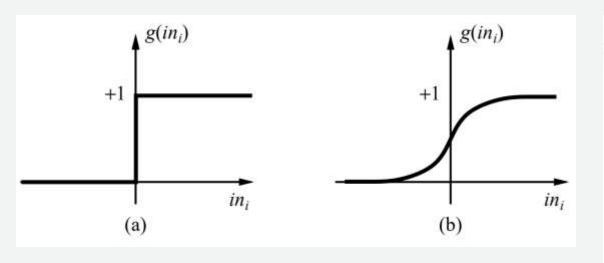
Neural Networks

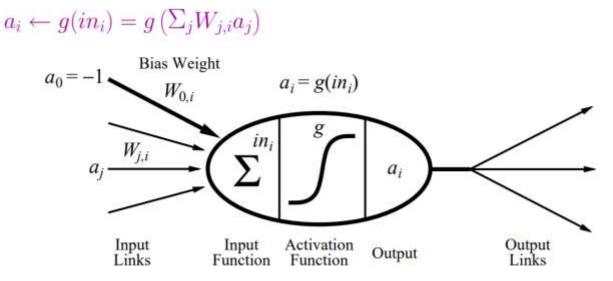
- Will focus mostly on algorithms which use combinations of neural networks
- Basic component of many important algorithms



Neural Networks: Neurons

- Output a squashed linear function of the inputs
- Activations can be step functions, inverse tangent functions, sigmoid functions, etc.

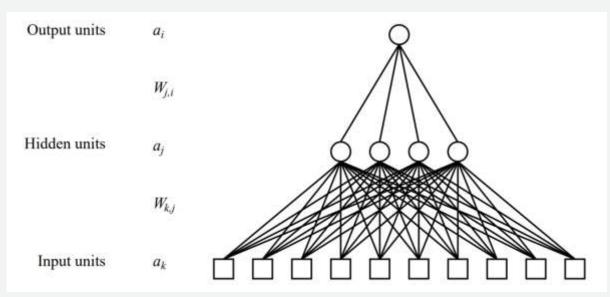




Neural Networks: Architecture

 Layers are fully connected, with # of hidden units and their layers chosen by hand

 Can have any number of input/output units

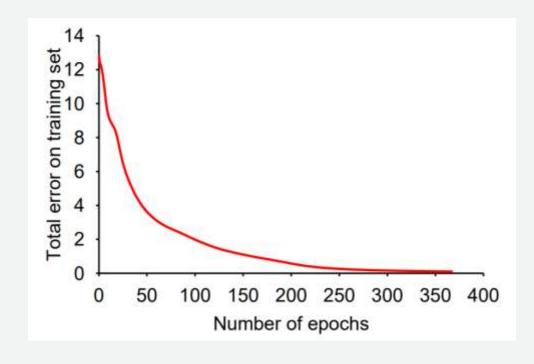


Neural Networks: Training

- From a theoretical perspective
 - NNs are usually trained by backpropagating (working back through the network and updating weight matricies W)
 - Gradients are automatically estimated by tensor libraries (tensorflow, pytorch, etc.)
 - At each epoch (instance of training on a batch of samples), the summed gradient updates for all examples

Neural Networks: Training

- Networks usually converge (stop learning)
 after many epochs (can vary a lot based
 on the model)
- Getting good results can be very very difficult
 - Often depends on the data!
 - Algorithms are fickle



Neural Networks: Training

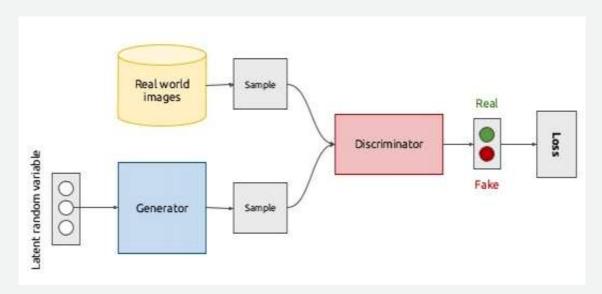
- Good resource for getting things to work (from Tesla Al director)
- Recipe (from karpathy):
 - 1. Spend LOTS of time visualizing the data
 - 2. Set up end-to-end training/eval skeleton
 - 3. Get a large enough model to overfit (be simple!)
 - 4. Regularize the model (prevent overfitting)
 - 5. Tune the model (hyperparameters!!)
 - 6. "Squeeze out the juice"

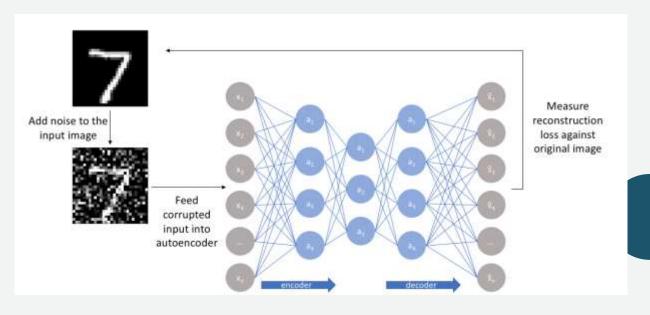
Hyperparameters

- Variable network parameters which change how well everything works
- NEED to optimize on these in order to be successful
- Examples:
 - Optimizer (thing which does SGD) and learning rate of optimizer
 - Regularization
 - # of nodes in each layer, # of layers, etc
 - Early stopping
 - Get more data!

Types of Models

- Deep networks
 - Very common, just means 'lots of layers'
- Generative Adversarial Networks (GAN)
 - Network that generates fake data competes with a network that discriminates fake vs. real
 - Gans work for both people and cats
- Autoencoder
 - Inputs match outputs
 - Learn some meaningful information about your data in an unsupervised fashion





Types of Models

- MANY more networks than we could possibly go through
- Let's build one!

